

House Sale Price Prediction

Alyssa Peterson

Sriram Ramadoss Venkata

Heather Simmons

Jessica Urban

Michael Xiong



Agenda

- Introduction
- About Dataset
- Linear Regression
- Neural Networks
- Random Forest
- Support Vector Machine
- Gaussian Mixture Model
- Algorithm Comparisons
- Q & A

Introduction

- Our goal for this project was to use regression and classification techniques in order to estimate the sale price of a house in King County, Washington given the feature and pricing data for around 21,000 houses sold within one year.

About Dataset



- Our dataset comes from a Kaggle competition.
- Our dataset contains house sale prices and its features for homes sold in King County, Washington between May 2014 and May 2015.
- King County is the most populous county in Washington and is included in the Seattle-Tacoma-Bellevue metropolitan statistical area. The county is considered the 13th most populous county in the United States.
- There are 21,613 observations in the dataset.
- There are 25 total attributes in the dataset, four of which we derived from current columns. We are planning to use 22 attributes in our models: all attributes except for date, latitude and longitude.

Dataset- Preprocessing

- During cleaning and preprocessing, we created four attributes derived from other attributes: Age, Age_revovated, Sqft_living15_diff, and Sqft_lot15_diff.
- We chose not to use the variable data, because it only shows us when the data was entered into the database.
- We chose not to use latitude and longitude because the attribute, Zipcode, contains the same information and was easier to work with in our models.
- We checked for missing variables and the dataset didn't contain any.
- We performed feature selection by looking at the correlation percentage of each attribute with price.

| Attribute | Percentage |
|--------------------|--------------|
| Bedrooms | 0.308349598 |
| Bathroom | 0.525137505 |
| Sqft_living | 0.702035055 |
| Sqft_lot | 0.089660861 |
| Floors | 0.256793888 |
| Waterfront | 0.266369434 |
| View | 0.397293488 |
| Condition | 0.036361789 |
| Grade | 0.667434256 |
| Sqft_above | 0.605567298 |
| Sqft_basement | 0.323816021 |
| Yr_built | 0.054011531 |
| Age | -0.054011531 |
| Yr_renovated | 0.126433793 |
| Age_renovation | -0.105754631 |
| Sqft_living15 | 0.585378904 |
| Sqft_living15_diff | 0.405391664 |
| Sqft_lot15 | 0.082447153 |
| Sqft_lot15_diff | 0.050590661 |

Outlier Detection using Cook's Distance

- We used Cook's distance for the original dataset and found one outlier. It was removed to create dataset 2.
- After applying Cook's distance to dataset 2, we found 608 outliers. They were removed to create dataset 3.
- The team used dataset 2 and dataset 3 for the project.

Linear Regression

- Used MATLAB to create 3 models
- Datasets : divided into 80% training, 20% test
 - Dataset 2 - one outlier removed
 - Dataset 3 – 608 outliers removed
- Algorithm
 - Model 1 used dataset 2
 - Model 2 used dataset 3 with 7 predictors removed
 - Model 3 used dataset 3 with 15 predictors removed
 - $\hat{y} = \text{predict}(\text{mdl}, X_{\text{new}})$

Linear Regression

Analytics Results:

Model 2 Training Metrics:

Root Mean Squared Error: \$96,800


R-squared: 0.879

Adjusted R-Squared 0.878

Model 2 Test Metrics:

Root Mean Squared Error: \$94,927.38

R-squared: 0.913223143



Neural Networks

- Used Matlab's built in functions: fitnet and feedforwardnet
- Tried two different methods: Levenberg-Marquardt and Bayesian
- Normalized with mapminmax to scale the targets, where the output of the network will be trained to produce outputs in the range $[-1, +1]$
- Started with 1 hidden layer and 10 neurons.
- Ran each combination again with 100 hidden neurons to determine which one would perform better.

Analytics Results:


Neural Networks

| 10 neurons | MSE | R-value | slope | y-intercept | MAE |
|--------------------------------|--------|---------|--------|-------------|--------|
| FeedForward Bayesian | 0.0019 | 0.8941 | 0.8185 | -0.1595 | 0.0289 |
| FitNet Bayesian | 0.0018 | 0.8951 | 0.8064 | -0.1701 | 0.0295 |
| FitNet LevenbergMarquardt | 0.0019 | 0.8931 | 0.797 | -0.1785 | 0.0294 |
| FeedForward LevenbergMarquardt | 0.0021 | 0.8826 | 0.7863 | -0.1865 | 0.0304 |
| | | | | | |
| 100 neurons | MSE | R-value | slope | y-intercept | MAE |
| FeedForward Bayesian | 0.0019 | 0.8983 | 0.8922 | -0.0944 | 0.025 |
| FitNet Bayesian | 0.0015 | 0.9142 | 0.8715 | -0.1128 | 0.0258 |
| FitNet LevenbergMarquardt | 0.0017 | 0.9009 | 0.8209 | -0.1574 | 0.0283 |
| FeedForward LevenbergMarquardt | 0.002 | 0.8878 | 0.8053 | -0.1714 | 0.03 |

Neural Networks

Analytics Results:

- **R-Squared score, Root Mean Squared Error (RMSE), Mean Absolute Error** of each model were calculated to find the quality and performance of the algorithm.
- **Fit Net Bayesian** with 100 hidden neurons performed the best
- Took nearly 2 hours to run the model
- R-Squared = 0.9142
- RMSE = 0.0015
- Mean Absolute Error = 0.0258




Random Forest

- Written in and executed in Python due to its extensive sklearn package.
- In this algorithm:
 - Transformed the target variable(price) to the log to correct for skew
 - Removed outliers using a function called median absolute deviation
 - Calculated the skew for each variable
 - Transformed variables if the skew was greater than 0.75
 - Applied a normalizer to each column
 - Ran Random Forest using 10 fold cross validation


Random Forest

- Feature Selection:
- Outputted feature importance coefficients, mapped them to their feature name, and sorted in decreasing order. Given our choice of model and methods for preprocessing data the most significant features are: 1. Grade 2. Zipcode 3. Sqft_living 4. Yr_built



Random Forest

- Analytics Results:
- **R-Squared score, Root Mean Squared Error (RMSE), Mean Absolute Error, and Explained Variance** of each model were calculated to find the quality and performance of the algorithm.
- **Model 1 performed the best**
- |R-Squared = 0.825 | Adj. R-Squared = 0.825| RMSE = 0.217 |Mean Absolute Error | = 0.157| Explained Variance Score= 0.825



SVM

- A **Support Vector Machine** (SVM) is a discriminative classifier formally defined by a separating hyperplane. Given labeled training data (supervised learning), the algorithm outputs an optimal hyperplane which categorizes new examples.

Advantages:

- SVMs produces large margin separating hyperplane, and efficient in higher dimension
- It maximizes the margin between points closest to the boundary
- SVMs only consider points near the margin (support vectors) – more robust

Disadvantages:

- Due to complexity of the algorithm it requires high amount of memory and takes long time to train the model and predict the test data
- The model is sensitive to optimal choice of kernel and regularization parameters

SVM

Data Preparation

- Used Correlation to figure out which predictor contribute more in prediction of prices
- Normalized all predictor to equal scale.
- Converted the target (Price – numerical data) to categorical values and into three bins.

Bin 1 0-300000

Bin2 300000-700000

Bin 3 700000+

SVM

Building Model

```
[rows useless] = size(X);  
G_Mat = zeros(rows, rows);  
rec = 0;  
for i = 1 : rows  
    d1 = X(i, :); % get one data point  
    rec = rec + 1; col = 0;  
    for j = 1 : rows  
        d2 = X(j, :); % get another data point  
        col = col + 1;  
        G_Mat(rec, col) = (1 + d1 * d2')^2; %  
    end  
end
```

Creating Gram matrix from training data

```
CVMD1 = fitcecoc(G_Mat, Y, 'ClassNames', classOrder);  
CMdl = CVMD1.Trained{1};
```

Fit the model using gram matrix

```
T_newX = [(1 + New_X * X').^2] ;  
[j, score] = predict(CVMD1, T_newX');
```

Predict the New X

SVM

Linear Support Vector Machine

| | Predicted – Class 1 | Predicted – Class 2 | Predicted – Class 3 | Total Actual |
|---------------------|------------------------|------------------------|------------------------|--------------|
| Actual – Class 1 | 0 | 0 | 0 | 0 |
| Actual – Class 2 | 4427 | 12582 | 4205 | 21214 |
| Actual – Class 3 | 10 | 182 | 207 | 399 |
| Total Predicted | 4437 | 12764 | 4412 | 21613 |

Accuracy = 0.58

Time taken = 10 mins

SVM

Kernel – Gaussian 14d

| | Predicted – Class 1 | Predicted – Class 2 | Predicted – Class 3 | Total Actual |
|---------------------|------------------------|------------------------|------------------------|--------------|
| Actual – Class 1 | 839 | 2197 | 330 | 3366 |
| Actual – Class 2 | 3510 | 6705 | 1880 | 12095 |
| Actual – Class 3 | 88 | 3142 | 2202 | 5432 |
| Total Predicted | 4437 | 12044 | 4412 | 20893 |

Accuracy = 0.46

Time taken = 1 hour

Kernel - Gaussian Infinite Dimension

SVM

| | Predicted – Class 1 | Predicted – Class 2 | Predicted – Class 3 | Total Actual |
|---------------------|------------------------|------------------------|------------------------|--------------|
| Actual – Class 1 | 1548 | 2881 | 8 | 4437 |
| Actual – Class 2 | 678 | 11473 | 613 | 12764 |
| Actual – Class 3 | 5 | 1572 | 2835 | 4412 |
| Total Predicted | 2231 | 15376 | 3456 | 21613 |



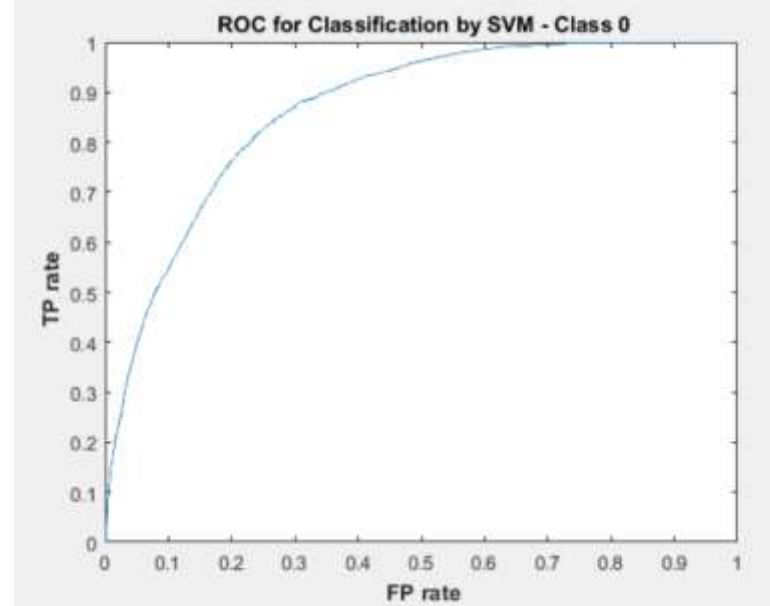
Accuracy = 0.73

Time taken = 2.5 hours

| | Class 1 | Class 2 | Class 3 |
|-----------|---------|---------|---------|
| Precision | 0.53 | 0.69 | 0.57 |
| Recall | 0.69 | 0.52 | 0.35 |
| AUC | 0.86 58 | 0.5426 | 0.20 |

Kernel - Gaussian Infinite Dimension

ROC



Gaussian Mixture Model

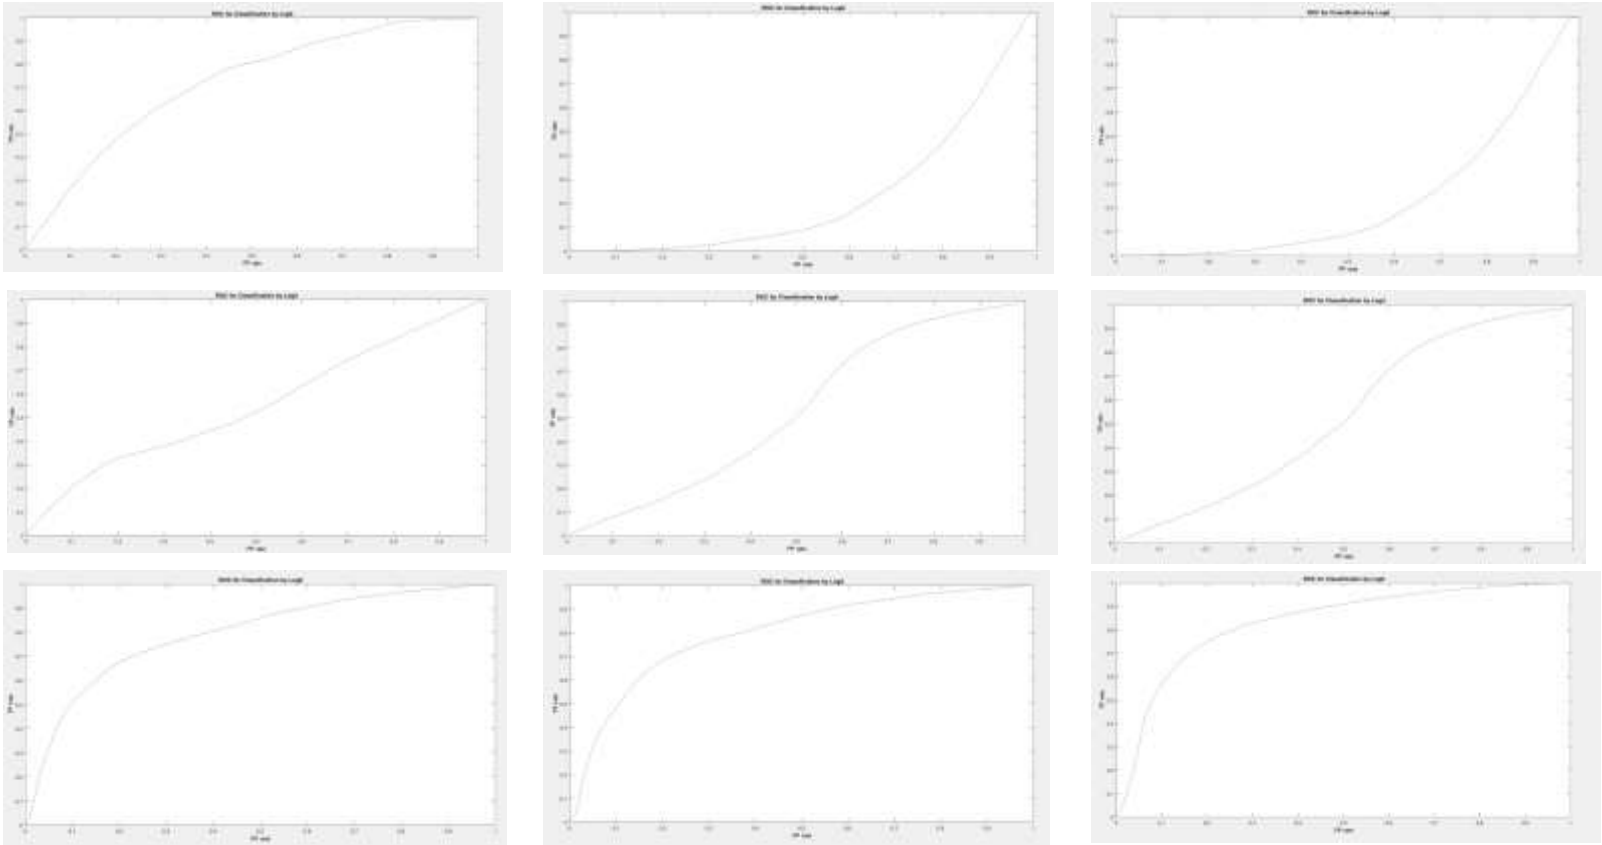
- Bucketed price values into three different groups:
 - Less than or equal to 300K (21.2%);
 - Greater than 300K
 - Less than 700K (58.4%); 700K or greater (20.4%).
- Principle Component Analysis was performed to select the best features.
 - Transformed data into a set of linearly uncorrelated variables.
 - Chose the two components that accounted for the majority of data variance.
- Ran with 1-3 clusters with different levels of regularization.
- Selected clusters with lowest Negative Log-Likelihood.
- Used posterior probability as a predictor to see if this improved the accuracy.
- The best overall model was with 3 clusters and 0 regularization.
- The best model accuracy and lowest negative log-likelihood was with the posterior probability with 0.05 regularization

Gaussian Mixture Model

Class 1:

Class 2:

Class 3:



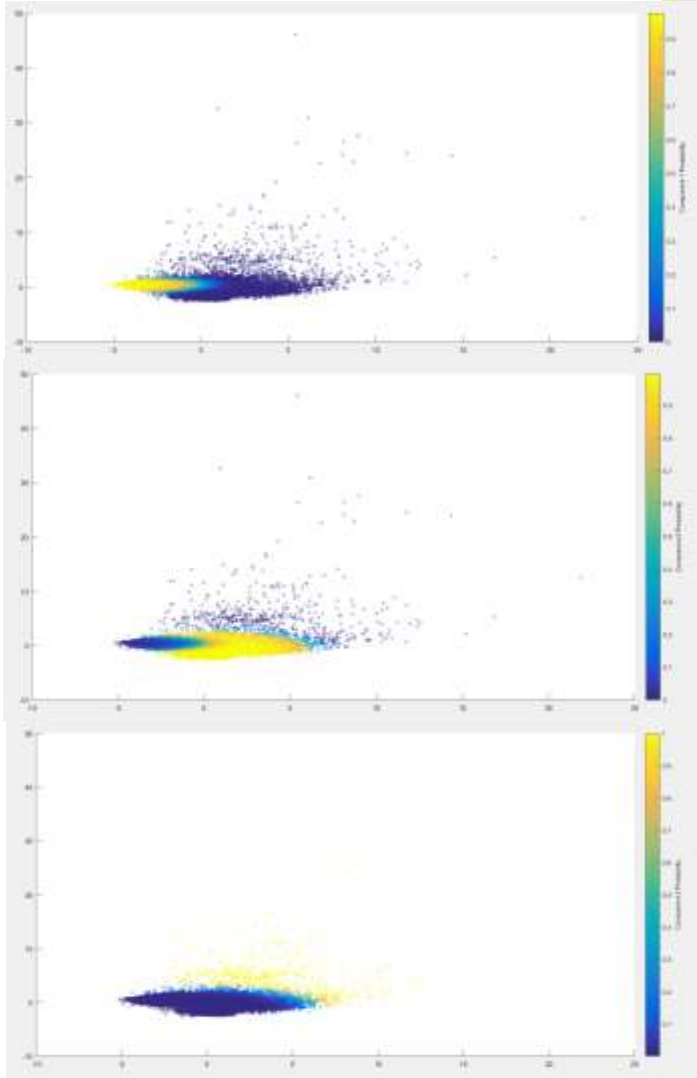
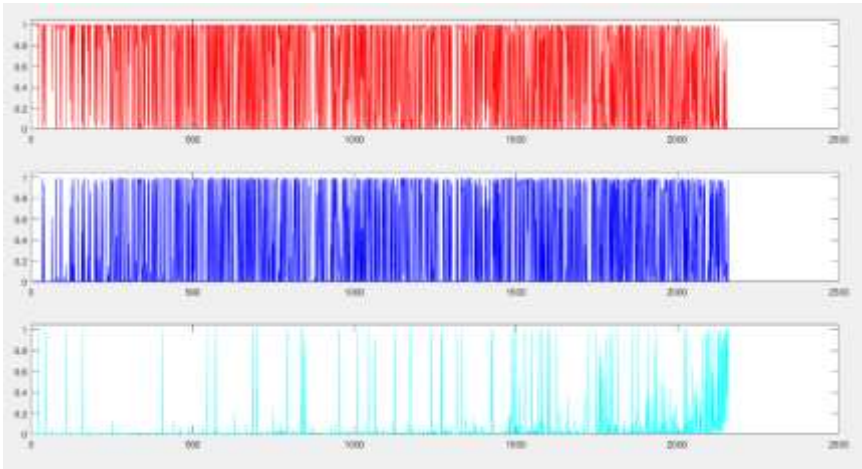

| | Model 1 | | | Model 2 | | | Model 3 | | |
|-----------|----------|--------|--------|----------|--------|--------|-----------|--------|--------|
| NLL | 6.85E+04 | | | 6.98E+04 | | | -6.25E+03 | | |
| Accuracy | 0.4928 | | | 0.5456 | | | 0.5483 | | |
| Class | One | Two | Three | One | Two | Three | One | Two | Three |
| Recall | 0.65 | 0.58 | 0.08 | 0.006 | 0.92 | 0.04 | 0.005 | 0.92 | 0.039 |
| Precision | 0.35 | 0.59 | 0.56 | 0.01 | 0.62 | 0.45 | 0.009 | 0.62 | 0.44 |
| AUC | 0.7176 | 0.5517 | 0.7943 | 0.2297 | 0.5316 | 0.8001 | 0.2337 | 0.5318 | 0.8339 |

Gaussian Mixture Model

Accuracy = 0.49
 Negative Log-Likelihood
 = 6.85e+04 = 68544

| CFM | Actual- Class 1 | Actual- Class 2 | Actual - Class 3 | Total Predic- ted |
|------------------------|--------------------|--------------------|---------------------|----------------------|
| Predicted - Class 1 | 3009 | 5086 | 618 | 8713 |
| Predicted- Class 2 | 1537 | 7284 | 3436 | 12257 |
| Predicted- Class 3 | 24 | 261 | 358 | 643 |
| Total Actual | 4570 | 12631 | 4412 | 21613 |

| | Class 1 | Class 2 | Class 3 |
|-----------|---------|---------|---------|
| Recall | 0.65 | 0.58 | 0.08 |
| Precision | 0.35 | 0.59 | 0.56 |
| AUC | 0.7176 | 0.5517 | 0.7943 |



Algorithm Comparisons

Regression Algorithms:

- Linear Regression:
 - R-squared: 0.913223143
 - Root Mean Squared Error: 0.349
- Neural Networks:
 - R-Squared = 0.9142
 - Root Mean Squared Error = 0.0015
 - Mean Absolute Error = 0.0015
- Random Forest:
 - R-squared = 0.825
 - Adjusted R-Squared = 0.825
 - Root Mean Squared Error = 0.217
 - Mean Absolute Error = 0.157

Algorithm Comparisons

Classification/Clustering Algorithms:

- Gaussian Mixture Model:

- Accuracy = 0.49
- Negative Log-Likelihood = $6.85e+04 = 68544$

| | Class 1 | Class 2 | Class 3 |
|-----------|---------|---------|---------|
| Recall | 0.65 | 0.58 | 0.08 |
| Precision | 0.35 | 0.59 | 0.56 |
| AUC | 0.7176 | 0.5517 | 0.7943 |

- Support Vector Machine:

- Accuracy= 0.73
- AUC :[0.86 58,0.5426,0.20]

| | Class 1 | Class 2 | Class 3 |
|-----------|---------|---------|---------|
| Precision | 0.5371 | 0.69 | 0.87 |
| Recall | 0.69 | 0.52 | 0.35 |
| AUC | 0.86 58 | 0.5426 | 0.20 |

Q & A